

AI TESTER SESSION REPORT — "THE COLLEGE BUDDIES RUN"

July 12, 2026 · Environment: Linux cloud sandbox against the live repo (PyQt6 offscreen; no game, no mic, no AI keys, no external network) · Both testers worked the same build, competitively.

MARCUS OKAFOR-REYES — "systems"

3

confirmed bugs · 25 checklist verdicts · ran all 18 official suites · 8 rendered-window evidence PNGs

DEVON "DEV" KOWALCZYK — "chaos" ★

WINNER

8

confirmed bugs · 18 checklist verdicts · 20 hostile item pastes · one very smug fake piper binary

THE ONE-PARAGRAPH VERDICT

The core of Kalandra is **solid**. The full official battery runs 290/290 on the stress suite; the PoB build pipeline survived a hand-crafted Witch export with decoy stats and garbage import codes; the dashboard constructs and renders end-to-end; config saves are genuinely atomic; secrets genuinely fail closed. The bugs the boys found cluster at the **edges**: two test suites that don't match the code they test, a filter editor that breaks its own byte-exact promise on Windows-style files, one crash on a None argument, a price-honesty loophole for headerless pastes, and four issues in the brand-new voice installer — **three of which were fixed and re-committed during this session** (regression-tested).

BUG SCOREBOARD BY SEVERITY

SEVERITY	COUNT	BUGS
CRASH	2	#4 install_voice --say traceback FIXED TODAY · #8 offline_craft_guidance(None) AttributeError
BROKEN	5	#1 db_status_checks dead-on-import · #2 clipboard bypasses provider seam · #5 stale-wav smoke test FIXED TODAY · #6 filter CRLF rewrite · #7 filter non-UTF8 mangling
CONFUSING	4	#3 TTS "ready" while broken · #9 config saved before smoke test FIXED TODAY · #10 headerless paste priced as currency · #11 OAuth "stored securely" when refused

WHAT THIS SESSION COULD AND COULDN'T COVER

Everything engine-level ran against your real code, including constructing the actual dashboard windows offscreen. What still needs a human at the Windows machine with PoE2 running: the mirror & medallion feel (A1–A12), in-game Ctrl+C popup focus behavior (E3/E4 live), ghost mode over the real game, voice in/out through real hardware (B-series live), the Orb's answer quality with a real AI key (stress rungs 1–5), and the four-authority accuracy checks against live PoB/poe.ninja/Craft of Exile/trade-site data. The printed checklist's Parts B and C remain the humans' job — the boys cleared the ground beneath them.

CONFIRMED BUG LOG

every bug reproduced twice · line-cited

#	WHERE	WHAT HAPPENS	EXPECTED	SEV / BY
1	tests/ db_status_checks.py:26	Imports <code>core_engine.database_handler.db_status</code> , which doesn't exist (logic lives in <code>scripts/db_status.py::gather()</code>). Whole suite dies before check #1 — 0/12.	Suite imports its subject and runs.	BROKEN M
2	gui_overlay/ mirror_window.py:3416	<code>on_clipboard_item</code> still calls <code>QApplication.clipboard()</code> + imports <code>parse_item_text</code> directly — bypassing the <code>item_in_hand</code> provider seam that W4-00 built and <code>provider_checks</code> enforces (46/49).	Clipboard flow goes through the provider registry like everything else.	BROKEN M
3	core_engine/ voice_engine.py:~62	<code>availability_message()</code> says "Local voice (TTS): ready" (find_spec only) while <code>pytttsx3.init()</code> actually raises and <code>list_voices()</code> returns nothing — two APIs in one module disagree.	"ready" means a voice can actually speak.	CONFUSING M
4	scripts/ install_voice.py --say	Stale <code>piper_path</code> in config → raw <code>FileNotFoundError</code> traceback from subprocess.	Clean "not installed — repair" message. FIXED + REGRESSION-TESTED TODAY	CRASH D
5	scripts/ install_voice.py smoke_test	Never deletes the previous <code>voice_test.wav</code> and ignores piper's exit code — piper can exit 1 writing nothing and the installer still reports "Voice test OK" off the stale file.	Test only passes on freshly produced audio. FIXED TODAY	BROKEN D
6	core_engine/ filter_editor.py load/ save	Every CRLF silently rewritten to LF on save (a flip + flip-back is NOT byte-identical: 84 → 77 bytes) — and this is a Windows game.	The documented "preserving every byte" round-trip.	BROKEN D
7	core_engine/ filter_editor.py:207	<code>errors='replace'</code> on load: non-UTF8 bytes (é as 0xE9) become U+FFFD and get written back — item names silently corrupted on save.	Refuse loudly or preserve original bytes.	BROKEN D
8	core_engine/ trade_tools.py:297	<code>offline_craft_guidance(None)</code> → <code>AttributeError</code> (line 292 shows None was meant to be tolerated).	Same graceful handling the empty-string goal gets.	CRASH D
9	scripts/ install_voice.py install()	Config keys + <code>voice_id="piper"</code> saved BEFORE the smoke test — a corrupt >1MB model leaves a broken voice selected.	Never auto-select a voice that failed its own test. FIXED TODAY	CONFUSING D
10	core_engine/ trade_tools.py:446	Paste with no <code>Item Class:</code> line but a currency-like name slips the honesty gate — a unique ring "priced" at ~180 ex.	Gear never gets an invented estimate, headerless or not.	CONFUSING D
11	gui_overlay/ mirror_window.py:1583	GGG OAuth branch ignores <code>set_secret()</code> 's return: says "Token stored securely in your OS keychain" even when keyring-less storage was refused; link silently gone next launch.	Check the return like the other save paths do.	CONFUSING D

ALSO FILED (NOT BUGS, BUT NOTED)

Non-English client pastes (German/Russian headers) parse without crashing but silently miss fields — worth a "unsupported client language" notice (Devon, E5). And the suite runner gap: 18 test files, no single command that runs them all and prints one scoreboard (Marcus's #1 wish).

CHECKLIST VERDICTS

✓ pass · ✗ fail · ? confusing · — needs a human / not reachable here

OFFICIAL TEST BATTERY (MARCUS)

18 suites run individually: **16 fully green** — stress_test 290/290, craft 111/111, companion 97/97, pob_sim 55/55, trade_query 35/35, nerf 43/43, tag_graph 28/28, scraper_tag 33/33 + 8 more. **Red:** db_status_checks 0/12 (bug #1), provider_checks 46/49 (bug #2).

ID	VERDICT	EVIDENCE (M = MARCUS, D = DEVON)
B1	✓ / ?	D: speak(""/None) returns False, on_done fires exactly once. M: availability_message truthful about missing key — but see bug #3 on "TTS ready".
B2	✓	M+D: list_voices() safe with/without piper; fake-piper render survived empty text, one emoji, 50k chars, null bytes; voice_rate clamps verified.
B3	✓	M: ai_respond with no key → graceful "[AI brain offline]" message, never an exception. D: hang containment — the 120s timeout is really passed to subprocess.
B4	✓	M: all 6 providers complete (label/service/kind/env/key_url/models). D: deleting the .onnx mid-session flips piper_available() and falls back cleanly.
B5– B7	—	Live AI answers, mic, audible speech: need keys/hardware. Humans' job (with the stress-rung page of the checklist).
C1– C6	✓ all	M: hand-built Witch/Infernalist PoB code round-trips — class/level/ascendancy, 3 skill groups ordered, actives/supports split, gear by slot incl. unique amulet, tree jewels, mainSocketGroup starring, DPS-max excludes decoy cost stat; 5 garbage codes → clean error dicts.
D1	✓	M: full KalandraDashboard constructs + renders offscreen (8 evidence PNGs).
D2– D4	✓	M: BuildTab constructs; scaling profile names fire/spell correctly on the test build; AI summary contains class/level/starred skill and respects budget.
D5– D6	—	PoB exe launch + pobb.in/ninja link resolution need Windows/network.
D7– D8	✓	M: IssueTracker + ReserveTracker persist across reload; reserve search works.
E1	✓	D: parser never raised across 20 hostile pastes (1 MB in 13 ms). M: offline data-layer suites all green.
E3– E5	? / ✓ / ?	D: honesty rule holds for real gear and currency stacks — except the headerless-paste loophole (bug #10); non-English pastes silently under-parse.
E7	✗	D: goal abuse fine (gibberish/essay/injection/100k chars) — but None crashes (bug #8).
E9	✗	D: byte-exact for LF files (incl. unicode, comments, 5000 blocks) — but CRLF rewrite + non-UTF8 mangling break the guarantee (bugs #6, #7). .bak is created.
F1– F2	✓	D: keyring removed → fails closed, no plaintext anywhere (repo-wide grep). GitHubSync tokenless → clean (False, message) everywhere.
F3– F5	✓	M: config atomic (tmp + os.replace, corrupt-save leaves old file intact); theme QSS parses; Kalandra Voice installer end-to-end with mocked downloads. D: offline modules degrade cleanly; oodle graceful on Linux; Corrupted Blood correction present and fires.
A1– A12	—	The mirror itself — drag, medallion clicks, ghost mode over the game — is physical-feel testing. Reserved for the humans with the printed checklist.

FAVORITE THING — "THE MOMENT IT FELT LIKE MAGIC"

MARCUS

The PoB bridge. I threw a hand-built Witch/Infernalist export at it — tree-socketed jewels, a decoy ManaCostPerSecondDPS stat, a disabled SRS group — and it got every single call right, including refusing to star the disabled group. That's the most load-bearing code in the app and it's the best code in the app.

DEVON

Honestly? The item parser. I fed it a megabyte of garbage, right-to-left Arabic, and null bytes, and it just... shrugged. 13 milliseconds. I was trying to make it cry and it made ME respect it. (Also my stale-wav zombie exploit, but that one's going in my highlight reel, not the app's.)

MOST ANNOYING THING — "WHAT YOU'D FIX FIRST"

MARCUS

The battery ships red out of the box: db_status_checks can't even import its subject, and provider_checks fails on code that's already merged. A test suite you can't trust green is worse than no suite — fix #1 and #2 before anything else.

DEVON

The filter editor swearing "we preserve every byte" in its docstring while quietly rewriting every CRLF in a Windows user's filter and lobotomizing any non-UTF8 item name. That's the one that eats a real player's loot filter someday.

ONE THING THEY'D ADD

MARCUS

A one-command runner: execute all tests/*.py, print one scoreboard. I looped 18 files by hand like an animal.

DEVON

A round-trip property test in CI: random byte-soup .filter files, load→save must equal input or refuse loudly. Same for parse_item_text with a no-crash guarantee. Two tests, permanent gremlin insurance.

THE BANTER (UNEDITED, AS CONTRACTUALLY REQUIRED)

MARCUS → DEVON

Devon can keep flailing at the fuzzer like a monkey with a keytar — chaos testing is just systems testing for people who can't read a call graph. Three confirmed, twice-reproduced, line-number-cited bugs, and I found one of them by noticing two APIs in the same module disagreeing with each other. Tell the gremlin that's what reading looks like.

DEVON → MARCUS

Hey spreadsheet boy — while you were color-coding your test matrix I got the installer to lie to my face with a wav file from the PAST, and I did it with a piper binary I wrote in 20 lines. That's 8 confirmed, all reproduced twice, zero network excuses. Bring more than a pivot table next time, Marcus.

RECOMMENDED FIX ORDER (JOINT, GRUDGINGLY AGREED UPON)

1. Bugs #1 + #2 — make the battery trustworthy-green again (small fixes, big trust). 2. Bugs #6 + #7 — filter editor byte-fidelity (CRLF + encoding); it touches player data. 3. Bug #8 — one-line None guard. 4. Bugs #10, #11, #3 — honesty polish (price gate, OAuth message, TTS "ready"). 5. Bugs #4/#5/#9 — **ALREADY FIXED & COMMITTED THIS SESSION**, regression-tested.

FINAL SCORE: DEVON 8 · MARCUS 3 — LOSER BUYS THE FIRST ROUND, PER THE RULES OF THE BET